
Title of the Thesis

Optional Subtitle of the Thesis

Author Full Name

Master in

Programme Name

Supervisor:

Prof. Supervisor Name, PhD,

Position, Department,

Iscte – Instituto Universitário de Lisboa

Co-Supervisor:

Prof. Co-Supervisor Name, PhD,

Position, Department,
Institution

Title of the Thesis

Optional Subtitle of the Thesis

Author Full Name

Master in

Programme Name

Supervisor:

Prof. Supervisor Name, PhD,

Position, Department,

Iscte – Instituto Universitário de Lisboa

Co-Supervisor:

Prof. Co-Supervisor Name, PhD,

Position, Department,
Institution

Month, Year

To my family.

Acknowledgments

I would like to thank my supervisor, **Prof. Supervisor Name**, for the guidance and support throughout this work.

Resumo

Este trabalho aborda o problema de ...

Palavras-chave: palavra-chave 1, palavra-chave 2, palavra-chave 3.

Abstract

This thesis addresses the problem of ...

Keywords: keyword 1, keyword 2, keyword 3.

Contents

Acknowledgments	i
Acknowledgments	i
Resumo	ii
Resumo	ii
Abstract	iii
Abstract	iii
List of Figures	vii
List of Tables	x
Chapter 1. Introduction	1
1.1. Motivation	1
1.2. Problem Statement	1
1.3. Research Questions	1
1.4. Objectives	1
1.5. Contributions	1
1.6. Document Structure	1
Chapter 2. Literature Review	2
2.1. Review Methodology	2
2.1.1. Search strategy	2
2.1.2. Inclusion and exclusion criteria	2
2.1.3. Results	3
2.2. Related Work	4
2.2.1. Cardiac MRI segmentation and datasets	4
2.2.2. Statistical shape models	5
2.2.3. Deep learning models for 3D shape reconstruction	7
2.2.4. Wall-thickness measurement techniques	9

2.2.5. Summary and research gap	10
Chapter 3. Methodology	12
3.1. Overview	12
3.2. Data Preparation	14
3.2.1. The missing ground-truth problem and motivation for SSM	15
3.2.2. Statistical Shape Model of the left ventricle	15
3.2.3. Synthetic end-diastolic dataset (SSM pre-training)	16
3.2.4. Real ED and ES labels: from SAX segmentation to 3D supervision	19
3.2.5. Shared cache representation and contour augmentation	25
3.2.6. Dataset composition and splits	27
3.3. Model Architecture	29
3.3.1. Contour encoder	30
3.3.2. Implicit decoder and positive-wall coupling	31
3.4. Training and Inference	32
3.4.1. Training data mix and augmentation	32
3.4.2. Loss function	33
3.4.3. Optimisation and hardware	34
3.4.4. Inference and analytic wall-thickness extraction	35
3.5. Wall-Thickness Evaluation Protocol	35
3.5.1. Selected thickness methods	36
3.5.2. Local three-dimensional wall thickness	39
3.5.3. Regional AHA-17 mapping	40
Chapter 4. Results and Evaluation	43
4.1. Experimental Setup	43
4.1.1. Metrics	43
4.1.2. Baselines	43
4.2. Results	44
4.2.1. Reconstruction Quality	44
4.2.2. Wall-Thickness Measurement	45
4.3. Discussion	45
4.4. Threats to Validity	46
Chapter 5. Conclusions	48

5.1. Summary	48
5.2. Revisiting the Research Questions	48
5.3. Contributions	48
5.4. Limitations	48
5.5. Future Work	48
References	49
References	49
Appendix A. Appendix Title	52
A.1. Section in the Appendix	52

List of Figures

2.1 PRISMA flow diagram summarising the literature selection process.	3
2.2 Overview of wall-thickness measurement approaches, ranging from simple distance queries to PDE-based transmural streamlines, with results aggregated into the standardised AHA 17-segment bullseye.	10
3.1 End-to-end overview of the thesis methodology. A representative SAX slice from ACDC patient002 illustrates the segmented clinical input and the corresponding endocardial, epicardial, and right-ventricular contours. The contour-extraction stage shows the sparse endocardial and epicardial observations used by CardioSDF, the mesh stage illustrates the reconstructed LV surface domain, and the final stages show local wall-thickness mapping and aggregation into the AHA-17 regional representation.	13
3.2 UK Digital Heart Project LV SSM: mean endocardial shape (left, with synthetic epicardium in orange) and the first three principal modes of variation shown at $\pm 3\sigma$. Mode 1 primarily controls overall cavity size; mode 2 modulates elongation along the long axis; mode 3 captures lateral asymmetry. The range of shape variation produced by these modes spans normal LV anatomy from compact to elongated morphologies.	16
3.3 Synthetic ED dataset generation pipeline. Panel (a) shows a scatter of accepted samples in the first two latent dimensions after Mahalanobis filtering. Panel (b) shows a generated endocardial mesh via Equation (3.1). Panel (c) illustrates the quality gate: accepted meshes (blue) are well within physiological bounds, while an outlier (red) is discarded. Panel (e) shows the sparse SAX contour rings extracted from the mesh at 10 axial levels. Panel (f) shows the 2048 query points stored in the occupancy cache, coloured by type (near-surface in blue, volumetric in grey).	18
3.4 Two-dimensional view of the clinical input. Representative SAX slices from ACDC patient002 at end-diastole (left, frame 1) and end-systole (right, frame 12), with the LV cavity, myocardium, and right-ventricle contours overlaid to confirm the label convention and frame selection.	20

- 3.5 Three-dimensional view of the same clinical input. The segmented ED SAX slices for ACDC patient002 are placed at their physical through-plane spacing, making the sparsity of the observation explicit: the input to the model is a stack of parallel contour rings, whereas the desired output is a continuous three-dimensional surface. 21
- 3.6 Real clinical label-to-mesh reconstruction, illustrated on ACDC patient002 at ED. Panel (a) shows the acquired SAX image with LV cavity (endocardial, blue), myocardium (epicardial boundary, orange), and RV (green) labels. Panel (b) converts these slice-wise labels to binary endocardial and epicardial masks. The pipeline then branches: panels (c)–(e) show the endocardial path, and panels (c′)–(e′) show the epicardial path. In each path, label contours are placed in physical 3D positions, a voxel surface makes discretisation artefacts visible, and Taubin smoothing illustrates their reduction. The figure is a visual explanation of the label-to-surface transition; the complete cache pipeline additionally performs the basal cap, fairing, decimation, quality checks, and mesh re-slicing described in Steps 5–6. 24
- 3.7 Contour-input augmentation applied to every training data stream. Each panel starts from the same ACDC patient002 ED contour stack; colours identify endocardial (blue) and epicardial (orange) points. The operations show representative draws of per-slice translation with point jitter, slice dropout, long-axis rotation with global scale jitter, and contour-point dropout. Only the encoder observation changes. The derived endocardial and epicardial meshes, supervision samples, query points, tissue labels, and phase label remain attached to the original case. 26
- 3.8 Number of candidate samples and how many survive the quality gates for each data stream. The real streams start from the 460 indexed patients; the synthetic stream starts from the 2048 sampled shape vectors. 28
- 3.9 Patient-level train, validation, and test split for the two real streams. Splitting by patient guarantees that no patient appears in more than one partition. 29
- 3.10 Conceptual comparison of the four wall-thickness estimators on a single myocardial cross-section, with the endocardium in blue and the epicardium in orange. (a) The EDT boundary sum adds the distances from an interior point to each boundary. (b) The Laplace field method follows the curved transmural streamlines of a harmonic potential, which cross the iso-potential surfaces (dashed lines) at right angles, and estimates thickness from the gradient magnitude. (c) The Yezzi–Prince method propagates two travel times u and v from opposite boundaries along that same curved field and sums

them. (d) The SDF cone-ray method casts a fan of rays within a cone around the endocardial normal and takes the median intersection distance (highlighted), making it robust to any single ill-oriented ray. 39

3.1 Local three-dimensional wall-thickness map on the reconstructed LV endocardial surface, shown from two opposing viewpoints. Each vertex is coloured by its local wall thickness on a fixed millimetre scale, so that the continuous spatial variation from the thin apex to the thicker base is visible directly on the geometry rather than only as regional averages. 40

3.1 Cutting the reconstructed LV surface into the AHA-17 segments. Panel (a) shows the endocardial surface partitioned along the long axis into basal, mid, apical, and apical-cap levels and angularly into sectors, with each of the 17 regions drawn in a distinct colour. Panel (b) unrolls the same colour-coded regions onto the standard AHA-17 bullseye, where the basal ring (segments 1–6) forms the outer annulus, the mid ring (7–12) the middle annulus, the apical ring (13–16) the inner annulus, and the apical cap (17) the central disc; the anatomical wall regions are labelled around the disc. 41

4.1 Mean wall thickness of the four selected methods on the reconstructed LV geometry, with whiskers spanning the 5th to 95th percentile. The dashed line marks the mean wall thickness measured directly from the input segmentation. 46

List of Tables

3.1 High-level summary of the CardioSDF methodology.	14
3.2 Configuration of the synthetic ED dataset generation pipeline.	18
3.3 Main operations in the real ED/ES mesh-preparation pipeline.	25
3.4 Summary of the three data streams prepared for CardioSDF.	27
3.5 Main fields in the shared training cache.	27
3.6 Final sample counts for the three data streams. Real streams report patients that pass the quality gates; the synthetic stream reports accepted meshes.	29
3.7 Main architecture settings of the model.	32
3.8 Purpose of the main training-loss terms.	34
3.9 Training configuration used for the model.	35
3.10 Wall-thickness methods selected for evaluation on the CardioSDF-generated LV geometry, with the Euclidean Distance Transform boundary sum kept only as a baseline.	37
4.1 Reconstruction quality of the final CardioSDF model on the test set. Chamfer distance and slice residual are in millimetres; the volume ratio is dimensionless with an ideal value of one.	44
4.2 Wall-thickness statistics for the four selected methods, computed on the CardioSDF-generated LV geometry. All values are in millimetres.	45

CHAPTER 1

Introduction

This chapter introduces the addressed problem, motivates the work, states the research questions, summarises the contributions, and outlines the remainder of the document.

1.1. Motivation

1.2. Problem Statement

1.3. Research Questions

This work addresses the following research questions:

RQ1.

RQ2.

RQ3.

1.4. Objectives

The main objectives are:

-
-

1.5. Contributions

The main contributions are:

-
-

1.6. Document Structure

The remainder of this document is organised as follows. Chapter 2 reviews the relevant literature and prior work. Chapter 3 describes the proposed approach. Chapter 4 presents and discusses the experimental results. Chapter 5 summarises the work and outlines directions for future work.

CHAPTER 2

Literature Review

This chapter presents the review methodology used to identify relevant literature and synthesises the related work on three-dimensional left-ventricle reconstruction from cardiac MRI and myocardial wall-thickness measurement.

2.1. Review Methodology

The literature search followed a structured approach to ensure broad coverage and relevance. This section describes the search strategy, selection criteria, and results.

2.1.1. Search strategy

The search was conducted across three major bibliographic databases: Scopus, PubMed, and IEEE Xplore, supplemented by Google Scholar for citation chaining. The temporal scope covered publications from 2000 to 2026, prioritising work from the last decade while retaining foundational methods essential for theoretical context.

Search terms were grouped into four thematic categories and combined using Boolean operators. The first category covered anatomy and imaging terms (“cardiac MRI”, “left ventricle”, “short-axis”, “SAX”, “myocardium”). The second addressed reconstruction (“3D reconstruction”, “mesh reconstruction”, “shape model”, “implicit surface”, “signed distance function”). The third targeted learning methods (“deep learning”, “graph neural network”, “point cloud”, “neural representation”). The fourth focused on measurement (“wall thickness”, “myocardial thickness”, “Laplace equation”, “transmural”).

The refined query applied consistently across databases was:

(“cardiac MRI” OR “left ventricle” OR “short-axis”) AND (“3D reconstruction” OR “statistical shape model” OR “signed distance function” OR “graph neural network” OR “implicit representation”) AND (“wall thickness” OR “mesh” OR “deep learning”)

2.1.2. Inclusion and exclusion criteria

Studies were included if they: (i) addressed three-dimensional reconstruction or shape modelling of the left ventricle, (ii) proposed or evaluated methods for myocardial wall-thickness




FIGURE 2.1. PRISMA flow diagram summarising the literature selection process.

measurement, (iii) developed geometric deep learning techniques applicable to anatomical meshes, or (iv) introduced benchmark cardiac MRI datasets used for LV analysis.

Studies were excluded if they: (i) focused exclusively on segmentation without reconstruction or shape analysis, (ii) addressed organs other than the heart without methodological generalisability, (iii) lacked peer review or sufficient methodological description, or (iv) were published before 2000 unless they constituted foundational work directly relevant to the methods reviewed.

2.1.3. Results

The initial search retrieved 142 records across the three databases. After removing 31 duplicates, 111 unique records were screened by title and abstract. Of these, 48 were selected for full-text assessment. Following detailed review against the inclusion criteria, 28 studies were retained. An additional 7 studies were identified through backward citation searching, yielding a final corpus of 35 publications.

The selected studies span five thematic areas: cardiac MRI acquisition and segmentation, statistical shape models, implicit neural representations, graph neural networks for mesh processing, and wall-thickness measurement. The following section synthesises the key findings from this corpus.

2.2. Related Work

This section synthesises the reviewed literature, organised by thematic area.

2.2.1. Cardiac MRI segmentation and datasets

Short-axis (SAX) cine MRI is the standard protocol for left-ventricle (LV) evaluation, producing 2D slices perpendicular to the long axis with high in-plane resolution (1–2 mm) but large inter-slice gaps (6–10 mm) (Cerqueira et al., 2002; Petitjean & Dacher, 2011). This anisotropic acquisition creates the fundamental challenge that motivates 3D reconstruction: the ventricle must be inferred volumetrically from sparse planar observations. Petitjean and Dacher (2011) reviewed over 100 segmentation methods and concluded that the inter-slice gap remains the single largest source of error in volumetric cardiac analysis, with apex and base delineation accounting for the majority of inter-observer variability.

The American Heart Association (AHA) 17-segment model, introduced by Cerqueira et al. (2002), standardised the partition of the myocardium into basal, mid-cavity, and apical regions. This model is universally adopted for reporting wall motion, perfusion, and thickness, and serves as the reference frame for all regional analysis discussed in subsequent sections.

Segmentation of endocardial and epicardial contours is a prerequisite for any downstream 3D analysis. Ronneberger et al. (2015) proposed U-Net, an encoder–decoder architecture with skip connections that was originally designed for neuronal structure segmentation but rapidly became the de facto backbone for cardiac MRI analysis. The architecture trains effectively from small annotated datasets and produces pixel-level predictions that, when applied to SAX slices, yield endocardial and epicardial contours with sub-pixel accuracy.

The Automated Cardiac Diagnosis Challenge (ACDC), organised by Bernard et al. (2018), provided the first standardised cardiac segmentation benchmark. The dataset comprises 100 patients distributed across five clinical groups: normal, dilated cardiomyopathy (DCM), hypertrophic cardiomyopathy (HCM), abnormal right ventricle, and previous myocardial infarction. Each case includes expert delineations at end-diastole (ED) and end-systole (ES). The challenge reported that top-performing deep learning methods achieved mean Dice scores exceeding 0.93

for the LV cavity and 0.89 for the myocardium, demonstrating that automated segmentation had reached expert-level performance for standard acquisitions.

Isensee et al. (2021) introduced nnU-Net, a self-configuring framework that automatically adapts network topology, patch size, preprocessing, and post-processing to each dataset. Applied to the ACDC benchmark, nnU-Net matched or exceeded all prior challenge submissions without any manual architecture design. The framework has since become the standard baseline against which new cardiac segmentation methods are evaluated.

The Multi-Centre, Multi-Vendor and Multi-Disease (M&Ms) challenge, described by Campello et al. (2021), extended evaluation to the clinically critical question of domain generalisation. The dataset includes 345 subjects acquired on scanners from four different vendors (Siemens, General Electric, Philips, and Canon) across six clinical centres. The challenge demonstrated that while methods trained on a single vendor achieved high accuracy on that vendor, performance dropped significantly when tested on unseen vendors — with Dice reductions of 5–15 percentage points in some cases. This finding motivated the development of domain-adaptation techniques and highlighted the importance of multi-source training data.

For population-level analysis, the UK Biobank (Sudlow et al., 2015) provides the largest available resource, with cardiac MRI acquired for over 40 000 participants using a standardised 1.5T protocol. Bai et al. (2018) developed a fully convolutional pipeline to segment this cohort at scale, producing automated delineations that, after quality control, serve as the foundation for large-scale shape modelling and epidemiological studies.

2.2.2. Statistical shape models

Statistical Shape Models (SSMs) encode anatomical variability by aligning a training population to point-to-point correspondence and applying PCA to extract the dominant modes of variation (Cootes et al., 1995; Heimann & Meinzer, 2009). Cootes et al. (1995) introduced Point Distribution Models, where a training set of shapes — each represented by k corresponding landmarks — is rigidly aligned via Procrustes analysis, averaged into a mean shape $\bar{x}$, and decomposed via PCA into principal modes Φ . A new shape is generated as $x = \bar{x} + \Phi b$, where b are shape parameters constrained within $\pm 3\sigma$ of each eigenvalue to remain anatomically plausible. In their experiments on hand and face contours, Cootes et al. (1995) showed that the first 5–10 modes captured over 95% of population variance, demonstrating the efficiency of the linear representation.

Heimann and Meinzer (2009) surveyed over 100 studies applying SSMs to medical image segmentation and identified the cardiac domain as particularly well-suited because of the

regularity of ventricular geometry and the availability of large training cohorts. Their review highlighted that SSM accuracy depends critically on three factors: the quality of inter-subject correspondence, the size and diversity of the training set, and the alignment method used.

Frangi et al. (2002) pioneered automatic construction of multi-object cardiac SSMs, proposing a non-rigid registration framework that establishes dense correspondence across subjects without manual landmarks. Applied to biventricular models, their approach captured coupled shape variation between the left and right ventricles, demonstrating that multi-structure models outperform single-structure ones for segmentation-by-fitting tasks.

Bai et al. (2015) constructed the most comprehensive cardiac atlas to date, using over 1000 high-resolution MRI scans from the UK Digital Heart Project. Their biventricular model captures both end-diastolic shape and wall motion through the cardiac cycle. The authors reported that the first 50 PCA modes explain approximately 90% of shape variance, and demonstrated that the atlas enables automatic segmentation of new subjects by fitting the model to image evidence with mean surface-to-surface errors below 1.5 mm.

de Marvao et al. (2014) applied the same UK Biobank-derived SSM to study myocardial hypertrophy and showed that atlas-based shape analysis detects subtle remodelling patterns that volumetric indices alone (such as LV mass) miss. Their study demonstrated substantially improved statistical power for detecting hypertrophic phenotypes when using shape modes compared to traditional clinical measurements, requiring fewer subjects to achieve equivalent sensitivity.

Medrano-Gracia et al. (2014) used a multi-ethnic cohort of over 1900 subjects to characterise normal LV shape variation. They reported that the dominant modes correspond to overall ventricular size, followed by eccentricity, apical shape, and regional wall curvature — confirming that cardiac shape cannot be adequately described by a single size parameter. Suinesiaputra et al. (2018) organised an international challenge to classify myocardial infarction from SSM-derived shape features, showing that machine learning classifiers trained on the first 20 shape modes achieved AUC values above 0.85 for distinguishing infarcted from healthy ventricles.

SSMs serve dual roles in reconstruction pipelines. First, they act as *geometric priors* that regularise shape recovery from sparse observations, constraining the solution to lie on or near the learned manifold. Second, they function as *synthetic data generators*: by sampling the latent space, one can produce thousands of anatomically plausible meshes for training data-hungry deep learning models without additional manual annotation (Heimann & Meinzer, 2009; Khalafvand et al., 2018). Khalafvand et al. (2018) demonstrated this second role by generating

SSM-sampled LV geometries to train computational fluid dynamics simulations, showing that the synthetic population reproduced the haemodynamic statistics of real subjects.

2.2.3. Deep learning models for 3D shape reconstruction

Classical reconstruction methods interpolate between segmented SAX slices using splines or deformable surfaces, but cannot recover geometric detail between slices, particularly near the apex and base (Heimann & Meinzer, 2009; Petitjean & Dacher, 2011). Petitjean and Dacher (2011) noted that linear interpolation between contours produces visible staircase artefacts and fails to reproduce the smooth taper of the LV apex, while parametric models (prolate spheroids, B-spline surfaces) impose smoothness but cannot express patient-specific morphological detail. Modern approaches replace interpolation with learned shape recovery, falling into two broad families: template deformation and implicit field prediction.

Template deformation methods maintain a fixed mesh topology and learn per-vertex displacements conditioned on input observations. Wang et al. (2018) proposed Pixel2Mesh, which takes a single RGB image as input and iteratively refines an ellipsoidal template mesh through three stages of graph convolutional layers. At each stage, vertex features are enriched with image features sampled via learned projections, and the mesh is refined through unpooling. The authors reported Chamfer distances below 0.5 on the ShapeNet benchmark, demonstrating that graph-based deformation can recover complex topology from minimal input. Choi et al. (2020) extended this paradigm with Pose2Mesh, which recovers full 3D human body meshes from sparse 2D pose keypoints. Their cascaded coarse-to-fine architecture first predicts a coarse mesh from joint positions and then refines vertex locations with finer graph convolutions, achieving state-of-the-art results on the Human3.6M benchmark. The architectural principle — recovering dense 3D geometry from sparse 2D observations — is directly analogous to the cardiac problem of inferring a full LV mesh from a handful of 2D SAX contours.

In the cardiac domain, template deformation preserves the topological consistency required for cross-subject comparison and downstream analysis such as wall-thickness computation (Bai et al., 2015; Suinesiaputra et al., 2014). Beetz et al. (2021) applied point-cloud variational autoencoders to biventricular shapes from the UK Biobank, learning a compact latent space that enables generation of subpopulation-specific anatomy models. Their decoder produces meshes with fixed vertex correspondence, making population statistics directly computable from the generated samples.

Implicit Neural Representations (INRs) take a different approach. Rather than deforming a template, they encode geometry as a continuous function whose zero-level set defines the surface. Park et al. (2019) introduced DeepSDF, which parameterises the signed distance function (SDF) of an entire shape class with a single neural network conditioned on a per-shape latent code. Given partial or noisy point-cloud observations, their auto-decoder architecture optimises the latent code at test time to find the shape in the learned space that best explains the input. On the ShapeNet benchmark, DeepSDF achieved reconstruction quality comparable to voxel-based methods while using an order of magnitude fewer parameters, and naturally supports shape interpolation in latent space. Mescheder et al. (2019) proposed Occupancy Networks, which predict binary inside/outside occupancy rather than signed distance. Their encoder processes a point cloud or image to produce a latent vector, and the decoder evaluates occupancy at arbitrary 3D query points. Both DeepSDF and Occupancy Networks produce resolution-independent surfaces extracted via the Marching Cubes algorithm (Lorensen & Cline, 1987), which generates a watertight triangulated mesh from the zero-crossing of the implicit field on a regular grid.

Two techniques are critical for applying INRs to thin cardiac structures where high-frequency geometric detail must be preserved. Tancik et al. (2020) demonstrated that standard MLPs exhibit spectral bias — they preferentially learn low-frequency components — and proposed mapping input coordinates through random Fourier features before the network. In their experiments on 2D image regression and 3D shape reconstruction, Fourier encoding substantially reduced reconstruction error compared to raw coordinate input, with the improvement most pronounced for high-frequency details such as thin structures and sharp edges. Gropp et al. (2020) introduced Implicit Geometric Regularisation, showing that enforcing the eikonal constraint ($\|\nabla f\| = 1$) as a loss term during training produces valid distance fields without requiring ground-truth SDF values at every point. Their approach enables training from raw point clouds alone, which is particularly relevant when dense SDF supervision is expensive to compute.

For encoding unordered input observations such as SAX contour points, the PointNet architecture (Qi et al., 2017) provides a permutation-invariant encoder via shared MLPs followed by symmetric max-pooling. Qi et al. (2017) demonstrated that this simple design achieves classification accuracy of 89.2% on the ModelNet40 benchmark and produces compact global descriptors that capture the essential geometry of the input point set regardless of its size or ordering. In cardiac applications, each SAX slice contributes a set of contour points that can be processed by a PointNet-style encoder to produce a latent representation of the observed geometry.

2.2.4. Wall-thickness measurement techniques

Myocardial wall thickness is a primary clinical biomarker for cardiac remodelling. Grossman et al. (1974) demonstrated that wall thickness is a principal determinant of left-ventricular diastolic stiffness, with hypertrophied ventricles exhibiting markedly elevated stiffness proportional to end-diastolic wall thickness. More recently, Huelnhagen et al. (2017) showed at 7.0T MRI that wall thickness correlates with tissue transverse relaxation properties, confirming its value as a non-invasive surrogate for myocardial tissue characterisation. The AHA 17-segment model (Cerqueira et al., 2002) provides the standard framework for reporting regional thickness values.

The accuracy of wall-thickness measurement depends critically on the quality of the underlying surface reconstruction, as any smoothing or geometric artefact propagates directly into the thickness estimate. The literature describes five families of measurement methods, each with distinct trade-offs.

The simplest approach uses nearest-neighbour distance: for each endocardial vertex, a KD-tree identifies the closest epicardial point, yielding a measurement in $O(N \log M)$ time. While computationally trivial (under 1 ms for typical cardiac meshes), this method provides only a lower bound on the true transmural thickness because it ignores the anatomical direction of measurement — the nearest point may lie tangentially rather than across the wall.

Ray-based methods address this limitation by casting a ray from each endocardial vertex along its surface normal until it intersects the epicardium. The intersection is computed via the Möller–Trumbore algorithm, accelerated with a Bounding Volume Hierarchy (BVH) to $O(\log F)$ per query. This approach measures transmurally but is sensitive to normal noise: a slightly rotated normal can produce a spurious long or short ray. The Shape Diameter Function variant improves robustness by casting a cone of K rays (typically $K=7$) at a half-angle of 30° around the normal and reporting the median intersection distance, effectively filtering outlier hits caused by local mesh irregularities.

Volumetric methods operate on voxelised representations. The Euclidean Distance Transform (EDT) computes, for each myocardial voxel, the distance to the nearest boundary, and the sum of the endocardial and epicardial distance fields approximates local thickness ($t = D_{\text{endo}} + D_{\text{epi}}$). This estimate is exact for planar parallel boundaries and overestimates by approximately 5–10% for the curvatures typical of the LV mid-wall. The method is computationally efficient ($O(V)$ via the Saito–Toriwaki algorithm) but limited by voxel resolution.

PDE-based methods represent the gold standard for transmural thickness. Jones et al. (2000) introduced the Laplacian approach — originally for cortical thickness in brain imaging — which

[Placeholder] Summary diagram of wall-thickness measurement families: (a) KD-tree nearest neighbour, (b) normal ray, (c) Laplace streamlines, (d) AHA-17 bullseye output.

FIGURE 2.2. Overview of wall-thickness measurement approaches, ranging from simple distance queries to PDE-based transmural streamlines, with results aggregated into the standardised AHA 17-segment bullseye.

solves $\nabla^2 \psi = 0$ with Dirichlet conditions ($\psi=0$ at the endocardium, $\psi=1$ at the epicardium). The uniqueness of the harmonic solution guarantees that the resulting streamlines never cross, producing a globally consistent transmural direction field. The local thickness is $t = 1/|\nabla \psi|$. While conceptually elegant, this formulation amplifies numerical errors where the gradient magnitude is small (thin walls, apex). Yezzi and Prince (2003) proposed an improvement: instead of dividing by $|\nabla \psi|$, they solve two advection equations ($\nabla \psi \cdot \nabla u = -1$ from the endocardium and $\nabla \psi \cdot \nabla v = -1$ from the epicardium), with the total thickness given directly as $t = u + v$. This Eulerian PDE formulation avoids the numerical instability of gradient division and produces identical results to the Laplacian method in smooth regions while being significantly more stable near the apex and in walls thinner than 4 mm.

Finally, mesh-based regularised correspondence methods combine an initial geometric estimate (normal projection or nearest neighbour) with iterative Laplacian smoothing on the mesh connectivity graph. At each iteration, each vertex's thickness is updated as a weighted average of its neighbours, anchored to the initial estimate. This produces spatially smooth thickness maps suitable for bullseye visualisation, at the cost of potentially attenuating real local thickness variation if the smoothing parameter is set too aggressively.

2.2.5. Summary and research gap

The reviewed literature shows that each component of the LV analysis pipeline is well-studied in isolation: segmentation is effectively solved by deep learning, SSMs encode population-level shape priors, implicit representations provide continuous surface modelling, and multiple wall-thickness algorithms exist. However, few studies integrate these components into a single pipeline for the specific setting addressed in this thesis. The missing combination is a method that encodes sparse SAX contour points, decodes paired endocardial and epicardial surfaces

as implicit signed distance fields, conditions the reconstruction on cardiac phase, guarantees positive wall thickness by construction, and trains jointly on real clinical data and SSM-derived synthetic samples. This gap motivates the approach described in the remaining chapters.

CHAPTER 3

Methodology

This chapter describes the methodology used to reconstruct 3D LV geometry from sparse short-axis contours and to measure myocardial wall thickness on the reconstructed model. The proposed pipeline, named CardioSDF, combines synthetic and real cardiac MRI-derived data, a phase-conditioned implicit neural representation, and a systematic comparison of local wall-thickness algorithms.

3.1. Overview

The complete thesis workflow can be read as a progression from sparse clinical observations to continuous three-dimensional geometry and then to quantitative wall-thickness summaries. It starts from segmented short-axis (SAX) cardiac MRI stacks or synthetic Statistical Shape Model (SSM) meshes. In both cases, the available geometry is converted into a common contour-based representation composed of endocardial and epicardial points sampled across SAX slices. These points are passed to a PointNet-style encoder that produces a compact latent code for the observed LV geometry. A signed-distance implicit decoder then predicts two coupled surfaces: the endocardium and the epicardium. The epicardial field is parameterised as a monotone offset from the endocardial field, which guarantees positive wall thickness by construction. The resulting surfaces are finally used to compute local thickness maps and regional AHA-17 summaries.

The methodological design is organised around four objectives. First, the input representation must be compatible with both synthetic and real cases. Second, the model must reconstruct smooth watertight LV surfaces from sparse 2D observations. Third, the decoder must produce a valid myocardial wall everywhere, not only at the observed slices. Fourth, the reconstructed geometry must support local wall-thickness measurements that can be visualised as surface colour maps and regional AHA-17 bullseye plots.

The chapter is therefore organised as a technical pipeline rather than as a single model description. Section 3.2 explains how synthetic and real data were converted into a shared cache representation. Section 3.3 describes the CardioSDF encoder and decoder, with particular attention to the monotone-epi parameterisation used to guarantee positive wall thickness. Section 3.4

presents the training and inference procedure, and Section 3.5 describes the wall-thickness algorithms used to evaluate local measurements on the reconstructed geometry.

Figure 3.1 summarises this end-to-end workflow and should be interpreted as an overview of the thesis methodology rather than as a figure belonging only to data preparation. The first panels illustrate the sparse SAX observation and contour extraction, the mesh panel represents the reconstructed LV surface domain, and the lower panels show how the reconstructed geometry is converted into local and regional wall-thickness measurements.




FIGURE 3.1. End-to-end overview of the thesis methodology. A representative SAX slice from ACDC patient002 illustrates the segmented clinical input and the corresponding endocardial, epicardial, and right-ventricular contours. The contour-extraction stage shows the sparse endocardial and epicardial observations used by CardioSDF, the mesh stage illustrates the reconstructed LV surface domain, and the final stages show local wall-thickness mapping and aggregation into the AHA-17 regional representation.

TABLE 3.1. High-level summary of the CardioSDF methodology.

Stage	Input / operation	Output
Data preparation	Synthetic SSM meshes and real ED/ES segmentations	Shared contour and query-point caches
Encoder	SAX contour points with tissue and phase labels	256-dimensional latent code z
Decoder	Latent code and Fourier-encoded query coordinates	Endocardial SDF, epicardial SDF, and bounded offset δ
Mesh extraction	Dense SDF samples on a 96^3 grid	Endocardial and epicardial triangular meshes
Thickness analysis	CardioSDF geometry and AHA orientation	Analytic thickness and selected local thickness maps

3.2. Data Preparation

A fundamental challenge in training CardioSDF is the absence of a paired dataset that provides both sparse SAX contour observations *and* corresponding validated three-dimensional LV meshes as ground truth. Clinical cardiac MRI datasets supply segmentation masks and image volumes, but they do not include the watertight endocardial and epicardial surfaces that an SDF decoder requires as supervision targets. This missing ground-truth problem is the starting point for the entire data-preparation strategy.

We addressed it with a two-stage approach. In the first stage, 1300 synthetic ED meshes were generated from the UK Digital Heart Project Statistical Shape Model (SSM) of the LV, providing anatomically plausible three-dimensional geometry with controlled shape variation. The model was pre-trained on this synthetic corpus, which established a geometric prior for LV shape variability. In the second stage, the training corpus was extended with meshes reconstructed directly from real clinical segmentations: ED and ES frames from the ACDC, M&Ms, and M&Ms-2 datasets. This fine-tuning stage introduced patient-specific anatomy, acquisition variability, and phase diversity that the SSM alone cannot provide.

All three data streams were converted into the same cache schema. The shared format allows the pre-training and fine-tuning stages to use the same training loop and loss function without any dataset-specific logic.

3.2.1. The missing ground-truth problem and motivation for SSM

The core task of this thesis is to train a neural network that predicts a continuous three-dimensional LV surface from sparse two-dimensional SAX contour observations. Such a network needs paired training examples of the form (sparse contours, 3D mesh). No publicly available cardiac MRI dataset supplies this pair directly. The standard cardiac MRI datasets such as ACDC (Bernard et al., 2018), M&Ms (Campello et al., 2021), and M&Ms-2 (Martín-Isla et al., 2023) provide per-slice segmentation masks, but not the watertight three-dimensional surfaces required by an SDF-based decoder.

Two complementary strategies were therefore adopted to create the missing supervision. First, a Statistical Shape Model was used to generate a large number of synthetic ED meshes that are anatomically valid by construction. This solved the ground-truth problem for the pre-training stage at the cost of not capturing real segmentation artefacts or scanner-specific anatomy. Second, real ED and ES segmentation masks were converted into watertight three-dimensional meshes through a controlled preprocessing pipeline. This provided the patient-specific geometry needed for fine-tuning, at the cost of requiring careful mesh reconstruction from inherently discrete voxel data.

The SSM pre-training approach is motivated by the observation that learning the geometry of the LV, rather than fitting individual segmentations, requires exposure to the full range of normal and pathological LV shapes. The SSM encodes this range compactly as a mean shape and a set of principal modes of variation, and sampling from its latent space produces diverse but anatomically constrained meshes. The real-data fine-tuning stage then ensures that the model generalises to the segmentation quality and acquisition conditions present in the evaluation datasets.

3.2.2. Statistical Shape Model of the left ventricle

A Statistical Shape Model represents the variability of an anatomical structure across a population by applying principal component analysis (PCA) to a set of aligned training meshes (Heimann & Meinzer, 2009). Each mesh is expressed as a vector of vertex coordinates; the mean of all training shapes forms the mean shape $\bar{X}$, and the eigenvectors of the covariance matrix of the centred shapes form the principal modes. Any shape in the population can be approximated as a linear combination of these modes, and new shapes that follow the observed statistics can be generated by sampling the latent coefficient vector.

This thesis uses the UK Digital Heart Project LV SSM (Bai et al., 2015), derived from UK Biobank cine MRI data of healthy volunteers. The model provides a mean endocardial surface

of 22 043 vertices together with 100 principal modes of variation stored as a $66\,129 \times 100$ matrix of mode vectors and a corresponding vector of 100 eigenvalues λ_i . The standard deviations $\sigma_i = \sqrt{\lambda_i}$ quantify the magnitude of each mode. A new endocardial mesh is generated by

$$X(b) = \bar{X} + \Phi b, \quad (3.1)$$

where $\bar{X}$ is the mean shape flattened to a vector, Φ is the matrix of mode vectors, and $b = [b_1, \dots, b_{100}]^\top$ is the sampled coefficient vector. The first few modes capture the largest sources of shape variation in the population; mode 1 primarily controls overall LV size, mode 2 encodes elongation along the long axis, and higher modes represent progressively more localised deformations, as illustrated in Figure 3.2.




FIGURE 3.2. UK Digital Heart Project LV SSM: mean endocardial shape (left, with synthetic epicardium in orange) and the first three principal modes of variation shown at $\pm 3\sigma$. Mode 1 primarily controls overall cavity size; mode 2 modulates elongation along the long axis; mode 3 captures lateral asymmetry. The range of shape variation produced by these modes spans normal LV anatomy from compact to elongated morphologies.

3.2.3. Synthetic end-diastolic dataset (SSM pre-training)

The synthetic ED dataset was built by sampling 1300 meshes from the SSM and converting each mesh into a training-compatible cache entry. The generation pipeline proceeds in five steps, illustrated in Figure 3.3.

Step 1: Latent coefficient sampling. Each candidate coefficient vector b was drawn by sampling each component as $b_i \sim \mathcal{N}(0, \sigma_i^2)$ and then applying two acceptance criteria:

$$\left| \frac{b_i}{\sigma_i} \right| \leq 3, \quad \sum_i \left(\frac{b_i}{\sigma_i} \right)^2 \leq \chi_{0.99, d}^2, \quad (3.2)$$

where σ_i is the standard deviation of mode i , d is the number of active modes, and $\chi_{0.99, d}^2$ is the 99th percentile of the chi-squared distribution with d degrees of freedom. The per-mode clipping prevents individual modes from dominating the shape, and the Mahalanobis bound prevents the combined deformation from falling outside the plausible population distribution. Accepted samples are scattered within a bounded ellipsoid in latent space, as shown in panel (a) of Figure 3.3.

Step 2: Mesh generation. Each accepted coefficient vector b is substituted into Equation (3.1) to produce a deformed endocardial vertex array, which is then triangulated using the face connectivity of the mean-shape mesh. The resulting triangular surface is watertight and oriented by construction.

Step 3: Quality filtering. Before a generated mesh is accepted into the training set, it is evaluated against four proxy metrics: enclosed volume (mL), surface area (cm²), sphericity index, and plausibility of the apical and basal normal directions. Samples that fall outside calibrated bounds for any of these metrics are discarded. Periodic exact audits using VTK mass-properties computation were also performed every 40 samples to verify the proxy metric calibration.

Step 4: Synthetic epicardium. The SSM provides only the endocardial surface. The epicardial surface was synthesised by offsetting each endocardial vertex p_i along its outward normal n_i by a spatially varying distance d_i ,

$$q_i = p_i + d_i n_i, \quad (3.3)$$

where d_i was set to approximately 10 mm near the base and 5 mm near the apex, with additive Gaussian noise of standard deviation 0.5 mm to avoid perfectly uniform geometry. This synthetic epicardium is not intended to reproduce clinical wall thickness but provides a consistent outer surface for SDF supervision during pre-training.

Step 5: SAX contour extraction and occupancy cache. Each pair of endocardial and epicardial surfaces was sliced at 10 evenly spaced axial levels to produce sparse SAX contour rings, matching the acquisition pattern of real cine MRI. For each case, 2048 query points were

sampled near the surfaces and in the surrounding volume, and signed-distance targets or inside/outside occupancy labels were computed and stored. These query points form the dense volumetric supervision used by the SDF loss terms during training.




FIGURE 3.3. Synthetic ED dataset generation pipeline. Panel (a) shows a scatter of accepted samples in the first two latent dimensions after Mahalanobis filtering. Panel (b) shows a generated endocardial mesh via Equation (3.1). Panel (c) illustrates the quality gate: accepted meshes (blue) are well within physiological bounds, while an outlier (red) is discarded. Panel (e) shows the sparse SAX contour rings extracted from the mesh at 10 axial levels. Panel (f) shows the 2048 query points stored in the occupancy cache, coloured by type (near-surface in blue, volumetric in grey).

TABLE 3.2. Configuration of the synthetic ED dataset generation pipeline.

Component	Configuration
Shape source	UK Digital Heart Project LV SSM, mean shape plus 100 PCA modes
Accepted meshes	1300 synthetic ED meshes
Latent constraints	Mahalanobis 99% bound and per-mode 3σ clipping
Quality checks	Volume, surface area, sphericity, and apex/base plausibility
Epicardium	Variable normal-offset surface, approximately 10 mm at base, 5 mm at apex
SAX extraction	10 SAX contour levels
Query cache	2048 near-surface and volumetric query points per case

3.2.4. Real ED and ES labels: from SAX segmentation to 3D supervision

SSM-derived meshes alone cannot capture the full diversity of clinical acquisition conditions: scanner differences, segmentation quality, image resolution, and the distinct geometry of the end-systolic phase. A second data stream was therefore created by reconstructing watertight three-dimensional LV meshes directly from the segmentation masks provided in the ACDC (Bernard et al., 2018), M&Ms (Campello et al., 2021), and M&Ms-2 (Martín-Isla et al., 2023) datasets, taking both the ED and ES frames so that the model is exposed to both cardiac phases.

These real ED and ES caches are used in the second optimisation stage, after the initial SSM-based training. The reason for this fine-tuning stage is that the synthetic corpus teaches a strong global LV-shape prior, but it does not fully reproduce clinical variability in contour quality, inter-dataset label style, and phase-specific contraction patterns. Fine-tuning on real ED and ES therefore reduces domain shift between synthetic supervision and the target clinical setting while preserving the geometric regularity learned from the SSM.

Operationally, the fine-tuning stage keeps the same CardioSDF architecture and the same multi-term loss from Section 3.4, but replaces the synthetic-only stream with real-derived cache samples from ED and ES frames. Because the real caches follow the same schema (contours, tissue labels, phase label, mesh supervision, and query points), the training loop can continue from the pre-trained checkpoint without changing model interfaces. The phase label then provides the conditioning signal that allows one shared model to adapt from diastolic to systolic geometry during continued training.

Before describing how these meshes are built, it is useful to look at the raw clinical input, because it makes clear why a reconstruction step is needed at all. Figure 3.4 shows two representative SAX slices from ACDC patient002, a dilated cardiomyopathy case whose metadata indicate ED frame 1 and ES frame 12 (Bernard et al., 2018), with the LV cavity, myocardium, and right-ventricular contours overlaid. Each frame is only a two-dimensional cross-section of the heart, and the segmentation is defined one slice at a time. Figure 3.5 stacks the segmented ED slices at their physical through-plane positions and exposes the central difficulty of the problem: a clinical acquisition provides only a small number of parallel SAX slices separated by several millimetres, so the observed geometry is a set of sparse contour rings rather than a continuous surface.




FIGURE 3.4. Two-dimensional view of the clinical input. Representative SAX slices from ACDC patient002 at end-diastole (left, frame 1) and end-systole (right, frame 12), with the LV cavity, myocardium, and right-ventricle contours overlaid to confirm the label convention and frame selection.




FIGURE 3.5. Three-dimensional view of the same clinical input. The segmented ED SAX slices for ACDC patient002 are placed at their physical through-plane spacing, making the sparsity of the observation explicit: the input to the model is a stack of parallel contour rings, whereas the desired output is a continuous three-dimensional surface.

The reconstruction pipeline closes this gap by converting each sparse label stack into a derived, watertight three-dimensional mesh. These meshes are *derived supervision*, not manually annotated 3D ground truth: their geometry is determined by the observed labels and by the explicitly stated resampling, smoothing, and capping operations. This distinction is important when interpreting the later reconstruction results. An earlier ES pipeline attempted to fit the UK Biobank SSM to systolic contours; this was abandoned because forcing a diastolic statistical prior onto systolic geometry produced anatomically unrealistic surfaces with collapsed walls. The direct reconstruction approach instead retains the patient-specific geometry present in the ED or ES labels. The two parallel endocardial and epicardial paths are shown in Figure 3.6; the six cache-construction steps are described below.

Step 1: NIfTI loading and RAS+ canonicalisation. Each segmentation volume is loaded from NIfTI format and reoriented to RAS+ canonical orientation using its affine. This places voxel axes in a consistent patient-coordinate convention while preserving physical spacing. It does not invent a new long axis; the acquired SAX stack supplies the through-plane direction, and an apex–base check is applied later before basal capping.

Step 2: Label harmonisation and binary mask construction. The ACDC, M&Ms, and M&Ms-2 datasets use compatible but not identical label encodings for the LV cavity, myocardium, and right ventricle. Each source is mapped to a common convention before the binary masks are formed. The endocardial mask is defined from the LV cavity label, and the epicardial mask from the union of LV cavity and myocardium labels:

$$M_{\text{endo}} = \mathbb{K}[L = L_{\text{LV}}], \quad M_{\text{epi}} = \mathbb{K}[L \in \{L_{\text{LV}}, L_{\text{MYO}}\}], \quad (3.4)$$

where L is the label image, L_{LV} is the LV cavity label, and L_{MYO} is the myocardium label.

Step 3: Isotropic resampling and label cleaning. The segmentation volume is resampled to a target spacing of 1.5 mm in all three dimensions using nearest-neighbour interpolation to avoid label blending. Within each SAX slice, the largest connected component of each label is retained and any holes are filled. This removes small disconnected label fragments that can arise near the basal and apical boundaries of the stack. The largest three-dimensional component is then retained for each mask, and the endocardial mask is unioned into the epicardial mask to enforce $M_{\text{endo}} \subseteq M_{\text{epi}}$. The apex is placed at low z and the base at high z using the relative cross-sectional areas at the two ends of the stack. This establishes a common direction for the basal operation without assuming that RAS+ alone defines cardiac anatomy.

Step 4: Marching cubes surface extraction. After three-dimensional component filtering, the binary mask is Gaussian-smoothed at a scale of 1.5 mm and its 0.5 isosurface is extracted with marching cubes (Lorensen & Cline, 1987). Smoothing before isosurface extraction suppresses the voxel staircase that would result from applying marching cubes directly to discrete labels. Panels (d) and (d') of Figure 3.6 provide an unsmoothed voxel-surface visualisation of this issue for the endocardial and epicardial masks, respectively.

Step 5: Post-processing—smoothing, basal cap, and decimation. The raw marching-cubes surface is processed in three sub-steps. Taubin λ/μ smoothing (Taubin, 1995) (25 iterations, $\lambda = 0.53$, $\mu = -0.55$) reduces residual staircase artefacts while approximately preserving global volume. Because the label stack ends at the valve plane, marching cubes can create

a rounded, non-anatomical basal dome. The upper 6% of the apex–base extent is therefore removed and the resulting opening is capped by a plane perpendicular to the established long-axis direction. Cotangent-Laplacian fairing (3 iterations) then polishes the cap join, and quadric-error decimation (Garland & Heckbert, 1997) limits the mesh to at most 4000 vertices. Panels (e) and (e′) of Figure 3.6 show the Taubin-smoothing effect; the cache meshes receive the complete post-processing sequence described here. The basal-capping operation removes the rounded dome that marching cubes creates at the valve plane using a plane perpendicular to the long axis, replacing it with an anatomically flat cap.

Step 6: Anatomical quality checks and contour extraction. Before a mesh is accepted into the training cache, seven anatomical prechecks are applied. Each criterion targets a distinct failure mode that the marching-cubes reconstruction can introduce when segmentation quality is low or when the label stack covers only part of the ventricle.

- (i) **Apical taper.** The cross-sectional area of the endocardial contour must decrease monotonically toward the apex. Cases where the apical end is wider than the mid-cavity are rejected, as this indicates an inverted or truncated stack.
- (ii) **Basal support.** The most basal slice must have a cross-sectional area consistent with the valve-plane region. Very small or absent basal contours suggest that the stack was acquired without reaching the base, which would make the basal cap geometrically unsupported.
- (iii) **Centreline drift.** The lateral displacement of slice-wise centroids along the long axis is bounded. Excessive drift indicates a pathological tilt or a segmentation artefact where contours shift discontinuously between adjacent slices.
- (iv) **Contour regularity.** Within each slice, the sampled endocardial and epicardial contour rings must be smooth and approximately convex. Jagged or self-intersecting contours produced by noisy segmentation boundaries are filtered out.
- (v) **Endocardium-inside-epicardium ordering.** For every axial level, the endocardial contour must lie inside the epicardial contour. Cases where label noise produces an inverted ordering are discarded to prevent the decoder from seeing anatomically impossible supervision.
- (vi) **Wall proxy thickness.** A slice-wise proximity check verifies that the mean distance between the endocardial and epicardial contours falls within a physiologically plausible range. Both paper-thin walls (segmentation bleed-through) and anomalously thick walls are rejected.

- (vii) **Long-axis ratio.** The ratio of the long-axis length to the maximum short-axis diameter must lie within the normal LV range. Cases with a ratio below the lower bound are prolate outliers, and those above the upper bound are either severely dilated or include spurious basal label fragments.

Cases that fail any of these checks are discarded from the cache. Accepted meshes are centred and normalised jointly, then intersected with 10 equally spaced axial planes after a 5% apical and basal margin. The largest intersection loop at each plane is resampled to 60 points for each available tissue. These are the contour points observed by the encoder. Separately, 2048 near-surface and volumetric query points are sampled and labelled by mesh occupancy for decoder supervision. Thus, the 2D-looking encoder input and the volumetric target are two representations of the same derived mesh, rather than independent annotations.




FIGURE 3.6. Real clinical label-to-mesh reconstruction, illustrated on ACDC patient002 at ED. Panel (a) shows the acquired SAX image with LV cavity (endocardial, blue), myocardium (epicardial boundary, orange), and RV (green) labels. Panel (b) converts these slice-wise labels to binary endocardial and epicardial masks. The pipeline then branches: panels (c)–(e) show the endocardial path, and panels (c')–(e') show the epicardial path. In each path, label contours are placed in physical 3D positions, a voxel surface makes discretisation artefacts visible, and Taubin smoothing illustrates their reduction. The figure is a visual explanation of the label-to-surface transition; the complete cache pipeline additionally performs the basal cap, fairing, decimation, quality checks, and mesh re-slicing described in Steps 5–6.

TABLE 3.3. Main operations in the real ED/ES mesh-preparation pipeline.

Operation	Purpose
RAS+ canonicalisation	Standardises patient-coordinate orientation while preserving physical spacing
Label harmonisation	Converts ACDC, M&Ms, and M&Ms-2 to a common LV/myocardium mask
Isotropic resampling at 1.5 mm	Reduces anisotropy before 3D surface extraction
Per-slice cleaning	Removes small 2D fragments and fills holes
3D largest-component filtering	Keeps the main anatomical LV region
Gaussian mask smoothing and marching cubes	Suppresses staircase artefacts and extracts the surface
Taubin smoothing and Laplacian fairing	Reduces residual artefacts while preserving global shape
Basal cap perpendicular to the apex–base axis	Replaces the non-anatomical basal dome at the valve-plane boundary
Quadric decimation to 4000 vertices	Produces a consistently sized mesh for all cases
Anatomical prechecks (7 criteria)	Rejects geometrically implausible or corrupted cases

3.2.5. Shared cache representation and contour augmentation

All prepared cases were stored in a shared format composed of contour points, tissue labels, cardiac phase labels, mesh-derived supervision points, and cached SDF or occupancy queries. The ED and ES caches use 10 SAX contour levels, 60 contour points per ring, and 2048 query points per case. This common schema underpins the two-stage training procedure. In the first stage, the model is trained on synthetic ED caches produced by the SSM, which establishes a global geometric prior. In the second stage, the model is fine-tuned on real ED and ES caches from ACDC, M&Ms, and M&Ms-2, which closes the gap between synthetic geometry and clinical anatomy. Because both stages share the same cache fields and the same training loop, the transition from pre-training to fine-tuning does not require any change to the model architecture or the loss function: the loader simply switches the data stream and training resumes from the pre-trained checkpoint.

The cache separates observation from supervision. The encoder observes only contour points, tissue identity, and phase. The decoder is supervised from surface samples, cached query points, and sign or SDF targets derived from the prepared meshes. This matches the intended use case: at inference time the model operates from sparse SAX observations alone.

Data augmentation is applied on the contour observation only, after the patient-level train/validation/test split. The mesh, sampled surfaces, query points, and occupancy or SDF targets remain fixed. This is deliberately an observation-robustness strategy, rather than a geometrically transformed copy of the anatomical target. Such label-preserving perturbations are commonly used to improve robustness when labelled medical data are limited (Chlap et al., 2021; Shorten & Khoshgoftaar, 2019).

Augmentation is applied during both training stages. In the pre-training stage it encourages the SSM-based encoder to tolerate contour variability. In the fine-tuning stage it prevents overfitting to the relatively small number of patient scans in the real cohort. The operations are identical in both stages; only the source of the contour data changes.

For every training sample, the loader draws small independent in-plane translations for each slice, point-wise in-plane jitter, global scale jitter, and contour-point dropout. Slice dropout and long-axis rotation are applied probabilistically, while retaining at least three slice levels. The endocardial and epicardial points of a given slice receive the same translation, preserving their local relationship. Tissue labels and the ED/ES phase label are not changed. As illustrated in Figure 3.7, the encoder is therefore trained to recover one fixed patient-specific target from plausible variations of its sparse SAX observation.




FIGURE 3.7. Contour-input augmentation applied to every training data stream.

Each panel starts from the same ACDC patient002 ED contour stack; colours identify endocardial (blue) and epicardial (orange) points. The operations show representative draws of per-slice translation with point jitter, slice dropout, long-axis rotation with global scale jitter, and contour-point dropout. Only the encoder observation changes. The derived endocardial and epicardial meshes, supervision samples, query points, tissue labels, and phase label remain attached to the original case.

Among these operations, the long-axis rotation rotates the whole contour stack rigidly about the ventricular long axis, so the relative arrangement of the SAX rings is preserved and only the observation frame changes.

TABLE 3.4. Summary of the three data streams prepared for CardioSDF.

Data stream	Phase	Source	Main role
Synthetic SSM	ED	UK Digital Heart Project LV SSM	Stage 1: SSM pre-training (geometric prior)
Real MRI	ED	ACDC, M&Ms, M&Ms-2	Stage 2: fine-tuning (clinical ED anatomy)
Real MRI	ES	ACDC, M&Ms, M&Ms-2	Stage 2: fine-tuning (phase conditioning, ES anatomy)

TABLE 3.5. Main fields in the shared training cache.

Field	Meaning	Source	Used by
Contour points	SAX endocardial and epicardial points	Synthetic slicing or real masks	Encoder and anchor loss
Tissue label	Endocardium or epicardium identifier	Contour extraction	Encoder tissue pooling
Phase label	ED or ES cardiac phase	Dataset metadata	Phase-conditioned encoder input
Surface samples	Points on endocardial and epicardial meshes	Mesh extraction	Surface and normal losses
Query points	Near-surface and volumetric samples	Occupancy/SDF cache	SDF, sign, and exterior losses
Mesh metadata	Vertices, faces, scale, and orientation	Preprocessing pipeline	Evaluation and visualisation

3.2.6. Dataset composition and splits

This subsection reports how many samples each data stream finally contributes and how the real streams are divided for training and evaluation. The real data come from two public collections: the Automated Cardiac Diagnosis Challenge (ACDC) and the second Multi-Centre,

Multi-Vendor and Multi-Disease Cardiac Segmentation challenge (M&Ms-2). Together they provide 460 patients (100 from ACDC and 360 from M&Ms-2). For each patient we reconstruct one mesh at end-diastole (ED, the fully filled ventricle) and one at end-systole (ES, the contracted ventricle), giving two real streams. The synthetic stream is generated from the UK Digital Heart Project SSM: of its 100 modes, the first 32 already capture 99% of the shape variance, so we sample within that sub-space. Drawing shape coefficients under the Mahalanobis bound of Equation (3.2) and rejecting implausible shapes yields 1300 accepted meshes out of 2048 candidates (63%).

Not every real patient produces a usable mesh. A segmentation may be a long-axis view, cover too few slices, or fail an anatomical plausibility check, such as an endocardium that is not fully enclosed by the epicardium. After these quality gates, 444 of the 460 patients yield a valid ED mesh and 423 yield a valid ES mesh. Figure 3.8 summarises this attrition for the three streams.




FIGURE 3.8. Number of candidate samples and how many survive the quality gates for each data stream. The real streams start from the 460 indexed patients; the synthetic stream starts from the 2048 sampled shape vectors.

To prevent information from leaking between training and evaluation, the real streams are split by patient: all meshes from a given patient stay within a single split. Figure 3.9 shows the resulting composition. The ED stream is divided into 310 training, 67 validation, and 67 test patients, and the ES stream into 338 training, 66 validation, and 19 test patients. The synthetic ED meshes are used only to enrich the pre-training stage and are never placed in the validation or test sets.




FIGURE 3.9. Patient-level train, validation, and test split for the two real streams. Splitting by patient guarantees that no patient appears in more than one partition.

Table 3.6 collects the final sample counts. In total, the model is trained and evaluated on 867 real LV meshes, supported by up to 1300 synthetic ED meshes during pre-training.

TABLE 3.6. Final sample counts for the three data streams. Real streams report patients that pass the quality gates; the synthetic stream reports accepted meshes.

Stream	Source	Phase	Valid meshes	Split		
				Train	Val.	Test
Real ED	ACDC + M&Ms-2	ED	444	310	67	67
Real ES	ACDC + M&Ms-2	ES	423	338	66	19
Synthetic ED	UK Biobank SSM	ED	1300	pre-training only		

3.3. Model Architecture

The model has two parts: an encoder and a decoder. The encoder reads the sparse contour points and summarises them into a single latent vector that describes the LV geometry of one case. The decoder takes that vector together with a 3D query point and returns how far the point is from the endocardial surface. Because the decoder can be evaluated at any point in space, the reconstructed surface is not tied to a fixed grid or a fixed number of vertices, and a mesh is produced only at the end by sampling the decoder densely. The encoder follows the point-cloud principle of Qi et al. (2017) and the implicit decoder follows the signed-distance idea of Park et al. (2019).

This split is a natural fit for the problem, because the input and the output have very different structure. The input is a small, unordered set of slice points, so a point-cloud encoder is appropriate: it does not require the points to be ordered or arranged on a grid. The output is a smooth, continuous surface, so an implicit decoder is appropriate: it can be sampled at any resolution and turned into a mesh only when needed. The key addition in this work is that the epicardium is not predicted on its own. Instead, it is defined as an outward offset from the endocardium, which forces the wall to have a positive thickness everywhere.

3.3.1. Contour encoder

The encoder receives, for each contour point, its 3D coordinates together with a label indicating whether the point is endocardial or epicardial and a label indicating the cardiac phase. A small shared network processes every point independently and turns it into a feature vector. The features of the endocardial points and the features of the epicardial points are then combined separately by taking, for each feature, its maximum over the points. This max operation is what makes the encoder independent of the order in which the points are listed, which is important because contour points have no natural ordering across patients or slices. The two resulting descriptors are joined and projected into a single 256-dimensional latent vector z .

More formally, let $P = \{p_i\}_{i=1}^N$ be the contour points of one case, where each p_i carries coordinates, a tissue label, and a phase label. The shared point network computes

$$h_i = \phi_{\text{enc}}(p_i), \quad (3.5)$$

where ϕ_{enc} is the shared network. The endocardial and epicardial descriptors are then obtained by max-pooling within each tissue,

$$z_{\text{endo}} = \max_{i \in \mathcal{E}} h_i, \quad z_{\text{epi}} = \max_{i \in \mathcal{P}} h_i, \quad (3.6)$$

where $\mathcal{E}$ and $\mathcal{P}$ are the endocardial and epicardial point sets, and the final latent code is

$$z = W_z[z_{\text{endo}}, z_{\text{epi}}] + b_z. \quad (3.7)$$

Pooling the two tissues separately keeps their information distinct. The endocardial points describe the shape of the cavity, while the epicardial points describe the outer boundary of the wall. If all points were pooled together, these two roles would be mixed and the decoder would have a harder time telling the cavity apart from the wall. Keeping them separate gives the latent code explicit access to both surfaces before they are combined.

3.3.2. Implicit decoder and positive-wall coupling

Before a query point is given to the decoder, its coordinates are expanded with a Fourier positional encoding. This simply adds a few sine and cosine copies of the coordinates at different frequencies, which helps the network represent fine geometric detail instead of only smooth, blurry shapes (Tancik et al., 2020). Three frequency bands are used. For a query point $x \in \mathbb{R}^3$ the encoding is

$$\gamma(x) = [x, \sin(2^0 \pi x), \cos(2^0 \pi x), \dots, \sin(2^{L-1} \pi x), \cos(2^{L-1} \pi x)], \quad (3.8)$$

with $L = 3$. The decoder itself is an eight-layer fully connected network of width 512 with a skip connection at the middle layer. It uses a smooth activation (softplus) so that the reconstructed surface is smooth and free of the sharp creases that a ReLU activation would introduce.

The decoder produces two outputs. The first is the signed distance to the endocardium, f_{endo} . The second is a positive offset δ , which represents the local wall thickness, that is, how far the epicardium lies outside the endocardium at that point. The offset is squeezed into a fixed positive range by a sigmoid,

$$\delta(x, z) = \tau_{\min} + (\delta_{\text{cap}} - \tau_{\min}) \sigma(g_{\delta}(x, z)), \quad (3.9)$$

where g_{δ} is the raw decoder output, σ is the sigmoid, and $\tau_{\min}$ and δ_{cap} are the smallest and largest allowed wall thickness. The epicardial field is then

$$f_{\text{epi}}(x, z) = f_{\text{endo}}(x, z) - \delta(x, z). \quad (3.10)$$

Because Equation (3.9) always keeps $\delta \geq \tau_{\min} > 0$, the epicardial surface can never touch or cross the endocardial surface: the two are always separated by at least $\tau_{\min}$. This is the mechanism that guarantees a positive wall thickness everywhere, rather than hoping the network learns it from data. In the final configuration $\tau_{\min} = 0.05$ and $\delta_{\text{cap}} = 0.45$ in normalised units, which corresponds to roughly 1.25–11.25 mm once mapped back to millimetres.

TABLE 3.7. Main architecture settings of the model.

Component	Configuration
Input features	3D coordinates, tissue label, and cardiac phase
Encoder type	Shared network with endocardial and epicardial max-pooling
Latent dimension	256
Positional encoding	Fourier features with $L = 3$ frequency bands
Decoder	8-layer network, width 512, skip connection at layer 4
Activation	Softplus (smooth)
Output heads	Endocardial signed distance and positive wall-thickness offset
Thickness range	$\tau_{\min} = 0.05$, $\delta_{\text{cap}} = 0.45$ in normalised units
Inference grid	96^3 query grid for marching cubes

3.4. Training and Inference

This section explains how the model is trained from the shared cache and how a surface mesh and a wall-thickness map are produced once training is finished. Training proceeds in two stages: the model is first pre-trained on the synthetic SSM cases and then fine-tuned on the real ED and ES cases. In both stages it must learn two things together: a valid distance field for the endocardium, and a sensible relationship between the endocardium and the epicardium.

3.4.1. Training data mix and augmentation

Training follows the two-stage strategy introduced in Section 3.2. In the first stage, the model is pre-trained only on the synthetic ED cases generated from the SSM. This large, controlled corpus teaches a strong prior for LV shape and tolerates aggressive augmentation. In the second stage, the pre-trained model is fine-tuned on the real cases reconstructed from ACDC, M&Ms, and M&Ms-2, covering both the ED and ES phases. Fine-tuning starts from the pre-trained weights and continues with the same architecture and loss, so the model keeps the geometric regularity learned from the SSM while adapting to real segmentation quality and to the systolic phase. Throughout both stages the phase label is provided as input, so a single model handles ED and ES, and augmentation is applied only to the contour input; the target meshes and the cached supervision are never changed.

The two stages are complementary. The synthetic ED cases give a large set of anatomically valid shapes but do not reproduce every artefact seen in real segmentations, which is why they are used for pre-training rather than as the final training data. The real ED and ES cases then supply genuine patient anatomy from several datasets and let the model learn how the shape changes between the two phases. To avoid information leaking between the training, validation, and test sets, the real cases are split by patient before any augmentation is applied. Because the augmentation only perturbs the observed contours, each real target stays tied to the anatomy of its own patient.

In practice, this means the encoder is trained to be robust to imperfect, noisy inputs, such as small shifts, jitter, missing slices, or dropped points, while the decoder is always supervised against a fixed, consistent 3D geometry.

3.4.2. Loss function

The training loss is a weighted sum of several terms, each responsible for one aspect of a good reconstruction. Some terms pull the predicted surfaces onto the target meshes, some keep the distance field well-behaved, some keep the surface faithful to the observed slices, and some prevent the wall from collapsing. The full objective is

$$\begin{aligned}\mathcal{L}_{\text{total}} = & \lambda_{\text{surf}}\mathcal{L}_{\text{surf}} + \lambda_{\text{eik}}\mathcal{L}_{\text{eik}} + \lambda_{\text{off}}\mathcal{L}_{\text{off}} + \lambda_{\text{normal}}\mathcal{L}_{\text{normal}} \\ & + \lambda_{\text{WT}}\mathcal{L}_{\text{WT}} + \lambda_{\text{L1}}\mathcal{L}_{\text{L1}} + \lambda_{\text{anchor}}\mathcal{L}_{\text{anchor}} + \lambda_{\text{sign}}\mathcal{L}_{\text{sign}} \\ & + \lambda_{\text{extsign}}\mathcal{L}_{\text{extsign}} + \lambda_{\text{bbox}}\mathcal{L}_{\text{bbox}} + \lambda_z\|z\|_2^2.\end{aligned}\tag{3.11}$$

The most important terms are the surface term, which places the two surfaces on their target meshes, and the eikonal term, which keeps the output a proper signed distance field by encouraging its gradient to have unit length (Gropp et al., 2020). The wall-thickness term reinforces the built-in lower bound on δ by discouraging the wall from ever thinning toward its minimum. The remaining terms handle finer effects such as surface orientation, agreement with the input contours, and the suppression of spurious surfaces outside the heart; their individual roles are listed in Table 3.8. The eikonal term is computed in full precision to keep its gradient calculation numerically stable.

TABLE 3.8. Purpose of the main training-loss terms.

Term	Purpose	Effect on reconstruction
$\mathcal{L}_{\text{surf}}$	Fits zero-level sets to surface samples	Places endocardial and epicardial surfaces on the target meshes
$\mathcal{L}_{\text{eik}}$	Encourages $\ \nabla f\ = 1$	Regularises the decoder as a signed-distance field
$\mathcal{L}_{\text{off}}$	Repels off-surface points from zero	Reduces spurious zero crossings away from the myocardium
$\mathcal{L}_{\text{normal}}$	Aligns predicted and target normals	Improves local surface orientation
$\mathcal{L}_{\text{WT}}$	Enforces the wall-thickness floor	Discourages collapse toward an unrealistically thin wall
$\mathcal{L}_{\text{L1}}$	Supervises cached query-point distances	Provides dense volumetric SDF information
$\mathcal{L}_{\text{anchor}}$	Anchors input contour points to the surface	Preserves agreement with observed SAX contours
Sign terms	Enforce inside/outside consistency	Stabilise the relative position of endocardium and epicardium
Bbox term	Pushes points outside the contour box outward	Suppresses exterior false surfaces

3.4.3. Optimisation and hardware

Both the SSM pre-training stage and the real-data fine-tuning stage use the same optimisation setup: the AdamW optimiser at a learning rate of 2×10^{-5} and weight decay 5×10^{-4} , with the learning rate following a cosine schedule and early stopping ending each stage once the validation loss stops improving. Fine-tuning simply resumes from the pre-trained weights rather than starting from scratch. To fit the model in the available GPU memory, training uses mixed precision, gradient clipping at 1.0 for stability, and gradient accumulation to reach a larger effective batch size. Each stage runs for at most 400 epochs, though early stopping usually ends it sooner. Training was carried out on two GPUs in parallel. A light weight-clipping strategy was used instead of the more expensive spectral normalisation, since the decoder is evaluated at many points per step and GPU memory is the main constraint.

TABLE 3.9. Training configuration used for the model.

Setting	Value
Training strategy	SSM pre-training, then fine-tuning on real ED/ES
Training mode	Combined ED and ES with phase conditioning
Augmentation	Contour-input perturbations for synthetic and real cases
Optimiser	AdamW
Learning rate	2×10^{-5}
Weight decay	5×10^{-4}
Maximum epochs	400 per stage
Batch size	8, with gradient accumulation
Gradient clipping	1.0
Precision	Automatic mixed precision
Hardware	2 GPUs in parallel
Inference grid	96^3

3.4.4. Inference and analytic wall-thickness extraction

At inference time, the model turns the input contours into the latent vector z and then evaluates the decoder on a regular 96^3 grid of points. Marching cubes reads the zero-level sets of f_{endo} and f_{epi} from this grid and produces the endocardial and epicardial meshes (Lorensen & Cline, 1987). No mesh repair or hole filling is applied afterwards; the surfaces come directly out of the model. A wall-thickness value is available for free at every endocardial vertex, because the decoder already predicts the offset δ there; converting it to millimetres gives

$$t_{\text{model}}(x) = s \delta(x, z), \quad (3.12)$$

where s is the scale that maps normalised coordinates back to millimetres. This thickness is not measured after the fact by comparing two meshes; it is a direct output of the model. The selected measurement algorithms in Section 3.5 are therefore used as independent points of comparison on the same geometry, not as the only way to obtain thickness.

3.5. Wall-Thickness Evaluation Protocol

The wall-thickness evaluation is performed on the CardioSDF-generated geometry. The segmentation volume is used to extract contour rings and to orient the AHA regional coordinate system, but the measured surfaces are the endocardial and epicardial meshes produced by the

trained model. This distinction is important: the evaluation tests thickness estimation on the reconstructed implicit model, not directly on raw segmentation meshes. Myocardial wall thickness is itself a routine indicator of cardiac health, because a wall that is abnormally thin can signal infarction or scarring, while an abnormally thick wall can signal hypertrophy. The evaluation is therefore organised in three steps: first the algorithms that measure thickness on the geometry, then the resulting local three-dimensional thickness field, and finally its aggregation into the clinical AHA-17 regional summary.

3.5.1. Selected thickness methods

An exploratory survey implemented ten candidate wall-thickness algorithms on the reconstructed geometry, spanning geometric distances, ray intersections, volumetric distance fields, graph distances, mesh regularisation, and partial-differential-equation (PDE) formulations. Reporting all ten would add redundancy without changing the conclusions, because several algorithms are minor variants of the same underlying principle. We therefore select three representative methods that together capture the accuracy–robustness trade-off relevant to a signed-distance reconstruction: the Laplace field method as the established transmural reference, the Yezzi–Prince Eulerian PDE method as a more recent and numerically more stable formulation, and the signed-distance cone-ray method as a promising surface-based estimator that aligns naturally with the SDF geometry produced by CardioSDF. The Euclidean Distance Transform (EDT) boundary sum is retained only as a volumetric baseline for context. Each selected method returns local thickness values in millimetres at endocardial vertices or projected endocardial locations. The outputs are summarised with mean, median, standard deviation, fifth percentile, ninety-fifth percentile, finite-value fraction, and runtime, and are visualised as 3D colour maps and AHA-17 bullseye plots.

TABLE 3.10. Wall-thickness methods selected for evaluation on the CardioSDF-generated LV geometry, with the Euclidean Distance Transform boundary sum kept only as a baseline.

Method	Category	Role	Main measurement principle
Laplace field	PDE	State of the art	Transmural field lines from a harmonic potential
Yezzi–Prince	PDE	Recent	Eulerian transport equations along the transmural field
SDF cone rays	Ray / SDF	Promising	Median of multiple cone rays around the surface normal
EDT boundary sum	Volumetric	Baseline	Sum of endocardial and epicardial distance transforms

As a volumetric baseline, the EDT boundary sum computes, at each myocardial voxel x , the sum of the distances to the endocardial and epicardial surfaces,

$$t^{\text{EDT}}(x) = D_{\text{endo}}(x) + D_{\text{epi}}(x), \quad (3.13)$$

where D_{endo} and D_{epi} are the Euclidean distance transforms of the two boundaries. This is exact for planar parallel walls but is limited by voxel resolution and slightly overestimates thickness in strongly curved regions, which motivates the three selected methods below.

The Laplace field method is the established transmural reference. It defines a smooth scalar field ψ over the myocardium by solving

$$\nabla^2 \psi = 0, \quad \psi|_{\text{endo}} = 0, \quad \psi|_{\text{epi}} = 1. \quad (3.14)$$

The local wall thickness is then estimated from the gradient magnitude along the transmural field. This family of methods is attractive because harmonic field lines do not cross, but numerical errors can appear where the wall is thin or where the field gradient is small (Jones et al., 2000). The more recent Yezzi–Prince approach addresses this by solving Eulerian transport equations along the same transmural field and combining the endocardial and epicardial travel times (Yezzi & Prince, 2003). In simplified form,

$$\nabla \psi \cdot \nabla u = -1, \quad t = u + v, \quad (3.15)$$

where u and v are the propagation distances from opposite boundaries.

The cone-ray method uses the endocardial surface normal but reduces sensitivity to a single noisy ray. Inspired by the shape diameter function (Shapira et al., 2008), for each endocardial point several rays are cast inside a cone around the normal direction, and the median valid hit distance is used:

$$t_i^{\text{cone}} = \text{median}_{k=1}^K d_{ik}, \quad K = 7. \quad (3.16)$$

In the notebook implementation the cone half-angle is 30° . This makes the estimate more robust than a single normal ray while still preserving an approximately transmural interpretation.

Figure 3.10 contrasts how the four estimators measure thickness on the same wall cross-section. The EDT baseline sums the distances from an interior point to both boundaries; the Laplace method follows the transmural streamlines of a harmonic potential whose nested iso-surfaces are shown as dashed lines; the Yezzi–Prince method integrates two travel times u and v along that same field and adds them; and the cone-ray method casts a fan of rays around the endocardial normal and keeps the median hit distance.




FIGURE 3.10. Conceptual comparison of the four wall-thickness estimators on a single myocardial cross-section, with the endocardium in blue and the epicardium in orange. (a) The EDT boundary sum adds the distances from an interior point to each boundary. (b) The Laplace field method follows the curved transmural streamlines of a harmonic potential, which cross the iso-potential surfaces (dashed lines) at right angles, and estimates thickness from the gradient magnitude. (c) The Yezzi–Prince method propagates two travel times u and v from opposite boundaries along that same curved field and sums them. (d) The SDF cone-ray method casts a fan of rays within a cone around the endocardial normal and takes the median intersection distance (highlighted), making it robust to any single ill-oriented ray.

3.5.2. Local three-dimensional wall thickness

Each of the selected methods produces, before any regional summary, a continuous thickness field defined over the whole reconstructed geometry. Because the CardioSDF decoder can be evaluated at any point in space, the endocardial and epicardial surfaces are watertight and densely sampled, and a thickness value is attached to each endocardial vertex. This produces a local three-dimensional wall-thickness map rather than a handful of per-slice numbers, as shown in Figure 3.11, where the LV surface is coloured by local thickness from the thin apical region to the thicker basal wall.

This local map is the primary output of the wall-thickness stage. Presenting the measurement first as a full surface field, and only afterwards as regional segment averages, makes the spatial distribution of thickness explicit: localised thinning or thickening remains visible on the surface even when it would be averaged away inside a single AHA segment. The colour map uses a fixed thickness range across cases so that maps from different patients and cardiac phases can be compared directly.




FIGURE 3.11. Local three-dimensional wall-thickness map on the reconstructed LV endocardial surface, shown from two opposing viewpoints. Each vertex is coloured by its local wall thickness on a fixed millimetre scale, so that the continuous spatial variation from the thin apex to the thicker base is visible directly on the geometry rather than only as regional averages.

3.5.3. Regional AHA-17 mapping

The dense surface field is finally aggregated into the 17-segment model of the American Heart Association (AHA), the accepted clinical language for reporting regional myocardial measurements (Cerqueira et al., 2002). Conventionally the wall thickness in each segment is read semi-manually from a few short-axis slices, which is observer-dependent, coarsely sampled, and sensitive to the choice of slice and measurement point. Here the same segment values are produced objectively from the local field of Section 3.5.2, by cutting the reconstructed surface into the AHA segments and averaging within each one.

The cut is defined in the LV's own anatomical frame, which first requires identifying the long-axis direction, distinguishing base from apex, and fixing a septal reference direction; the

wall-thickness notebook includes an anatomical QA step that visualises this base/apex orientation and septal alignment before the segments are formed. The surface is then partitioned in two stages, as illustrated in Figure 3.12. First, along the long axis, the endocardial vertices are split by their normalised apex–base position into a basal ring, a mid-cavity ring, an apical ring, and a small apical cap. Second, within each ring the vertices are divided angularly around the long-axis centreline: the basal and mid rings are cut into six sectors each and the apical ring into four, while the apical cap forms the single seventeenth segment. Numbering the sectors from the anterior–septal reference and unrolling the rings onto concentric circles yields the familiar bullseye, in which the basal ring is the outer annulus, the mid ring the middle annulus, the apical ring the inner annulus, and the apical cap the central disc.




FIGURE 3.12. Cutting the reconstructed LV surface into the AHA-17 segments. Panel (a) shows the endocardial surface partitioned along the long axis into basal, mid, apical, and apical-cap levels and angularly into sectors, with each of the 17 regions drawn in a distinct colour. Panel (b) unrolls the same colour-coded regions onto the standard AHA-17 bullseye, where the basal ring (segments 1–6) forms the outer annulus, the mid ring (7–12) the middle annulus, the apical ring (13–16) the inner annulus, and the apical cap (17) the central disc; the anatomical wall regions are labelled around the disc.

The AHA aggregation does not replace the per-vertex thickness maps; it provides a clinically familiar regional summary of the same local measurements. For each method, the thickness values of the vertices falling in a segment are combined using robust summaries such as the mean and percentiles, so that local outliers do not dominate the regional interpretation.

This evaluation should be interpreted as an algorithmic comparison and reproducibility study rather than as clinical validation. Unless expert manual wall-thickness measurements are added, the thesis can conclude that CardioSDF provides an objective and locally resolved computational framework, but it should not claim replacement of clinical measurement protocols.

Results and Evaluation

This chapter presents the experimental evaluation of the approach described in Chapter 3.

4.1. Experimental Setup

We evaluate the final CardioSDF model, obtained by pre-training on the synthetic SSM meshes and then fine-tuning on the real end-diastolic (ED) and end-systolic (ES) data, as described in Chapter 3. The datasets, quality gates, and patient-level splits are reported in Section 3.2.6; all numbers below are computed on the held-out test partitions, which contain no patient seen during training. Two questions guide the evaluation: how faithfully the model reconstructs the LV surfaces from sparse contours, and how the analytic wall thickness it produces compares with independent measurement algorithms applied to the same geometry.

4.1.1. Metrics

Reconstruction quality is measured with four quantities. The *Chamfer distance* averages, for every reconstructed vertex, the distance to the nearest ground-truth surface point and vice versa, reported separately for the endocardium and the epicardium in millimetres; smaller is better. The *slice residual* is the mean absolute signed-distance value at the input contour points, and measures how closely the reconstructed surface passes through the observed short-axis (SAX) rings, again in millimetres. The *watertight rate* is the fraction of reconstructed meshes that form a closed surface with no holes. The *volume ratio* is the reconstructed cavity or wall volume divided by its ground-truth value, so that a value of one indicates no systematic over- or under-estimation.

Wall thickness is summarised by its mean, its dispersion (standard deviation), and its 5th and 95th percentiles across the myocardial surface, all in millimetres. These percentiles describe the thin apical and thick basal extremes without being distorted by isolated outliers.

4.1.2. Baselines

Because no paired ground-truth 3D mesh exists for the clinical scans (see Section 3.2.1), the reconstruction is compared against the derived meshes built from the segmentations, and the

analytic wall thickness is compared against the four measurement algorithms selected in Section 3.5.1: the Laplace field method, the Yezzi–Prince method, the SDF cone-ray method, and the Euclidean Distance Transform (EDT) boundary sum kept as a volumetric baseline. As an external plausibility anchor, we also report the mean wall thickness obtained directly from the input segmentation.

4.2. Results

This section reports first the geometric quality of the reconstruction and then the wall-thickness measurements obtained from the reconstructed geometry.

4.2.1. Reconstruction Quality

Table 4.1 summarises the reconstruction accuracy on the test set. The endocardial and epicardial surfaces are recovered with a mean Chamfer distance of 1.12 mm and 1.17 mm respectively, which is close to the 1.0 mm isotropic resolution of the prepared data and therefore near the limit of what the derived meshes can resolve. The mean slice residual of 0.57 mm shows that the reconstructed surfaces pass through the observed SAX contours with sub-millimetre agreement, confirming that the model honours its input. The cavity and wall volumes are reproduced almost exactly, with volume ratios of 1.00 for the endocardium and 1.02 for the epicardium.

Every reconstructed mesh is watertight, for both the endocardium and the epicardium. This is a direct consequence of reading the surfaces as zero-level sets of a continuous signed-distance field rather than assembling them from per-slice contours, and it holds without any mesh repair, hole filling, or cap synthesis. The monotone-epi coupling also guarantees a strictly positive wall: the analytic thickness never collapses to zero, with a 5th percentile of 2.33 mm that stays well above the enforced floor.

TABLE 4.1. Reconstruction quality of the final CardioSDF model on the test set. Chamfer distance and slice residual are in millimetres; the volume ratio is dimensionless with an ideal value of one.

Metric	Value
Endocardium Chamfer distance (mm)	1.12
Epicardium Chamfer distance (mm)	1.17
Slice residual (mm)	0.57
Endocardium volume ratio	1.00
Epicardium volume ratio	1.02
Watertight rate (endocardium)	100%
Watertight rate (epicardium)	100%

4.2.2. Wall-Thickness Measurement

The analytic wall thickness read directly from the decoder has a mean of 5.31 ± 0.41 mm across the test set, with a 5th percentile of 2.33 mm and a 95th percentile of 6.67 mm. These values fall within the normal range for the LV myocardium and are consistent with the four independent algorithms applied to the same reconstructed surfaces.

Table 4.2 lists the mean, standard deviation, and 5th and 95th percentiles for each of the four selected methods, and Figure 4.1 shows the same comparison graphically. The three transmural estimators agree to within roughly one millimetre of the mean: the Laplace field reports 5.08 mm, the SDF cone-ray method 5.47 mm, and the Yezzi–Prince method 4.30 mm. The EDT boundary sum returns a lower mean of 3.58 mm, as expected for a voxel-based estimator that measures the shortest transmural distance and therefore reads low in curved regions; it is retained only as a baseline. For reference, the mean wall thickness measured directly from the input segmentation is 3.61 mm, and a plain nearest-neighbour measurement on the model surface gives 4.55 mm, so the surface-based and field-based estimators bracket the segmentation reference from above, as expected when the reconstruction fills in a smooth, complete wall rather than a slice-limited one.

TABLE 4.2. Wall-thickness statistics for the four selected methods, computed on the CardioSDF-generated LV geometry. All values are in millimetres.

Method	Mean	Std.	p_5	p_{95}
Laplace field (state of the art)	5.08	1.70	2.92	8.62
Yezzi–Prince (recent)	4.30	1.43	0.96	5.80
SDF cone rays (promising)	5.47	1.15	3.25	7.16
EDT boundary sum (baseline)	3.58	0.86	1.49	4.81

4.3. Discussion

The results support the two central design choices of CardioSDF. First, encoding the LV as a signed-distance field yields watertight surfaces in every test case without any post-processing, which removes the mesh-repair heuristics that a contour-stacking pipeline would otherwise require. Second, the monotone-epi coupling delivers a wall that is positive everywhere, so the analytic thickness is usable as a measurement rather than merely a by-product: its mean of 5.31 mm and its percentile range agree with the established Laplace field and with the SDF cone-ray estimator to within about one millimetre.

The spread between methods is informative rather than contradictory. Field-based transmural estimators and surface-based ray estimators measure thickness along slightly different




FIGURE 4.1. Mean wall thickness of the four selected methods on the reconstructed LV geometry, with whiskers spanning the 5th to 95th percentile. The dashed line marks the mean wall thickness measured directly from the input segmentation.

paths, and the EDT baseline reads systematically lower because straight-line voxel distances underestimate thickness where the wall curves. The agreement of the three transmural methods, together with the sub-millimetre slice residual and near-unit volume ratios, indicates that the reconstructed geometry is both accurate at the observed slices and plausible between them, which is precisely where sparse SAX acquisitions provide no direct evidence.

4.4. Threats to Validity

The most important limitation concerns the ground truth. The reconstruction is evaluated against meshes *derived* from the segmentations, not against independently acquired 3D surfaces, so the Chamfer distances measure agreement with a derived target rather than with true anatomy. For the same reason, the wall-thickness comparison is an inter-method and reproducibility study on shared geometry, not a clinical validation: without expert manual wall-thickness measurements, the thesis can claim an objective and locally resolved computational framework, but not the replacement of clinical measurement protocols.

Several further factors bound the external validity of the results. The real data come from two public collections (ACDC and M&Ms-2), and the ES test partition is small, so per-phase estimates carry more uncertainty than the ED ones. The wall-thickness percentiles depend on the AHA orientation and on the 96^3 inference grid, both of which introduce discretisation effects that a finer grid or an alternative anatomical frame could shift. Finally, the synthetic epicardium

used during pre-training is a geometric construction rather than a measured surface, which the fine-tuning stage on real data is designed to correct but cannot fully rule out as a residual prior.

CHAPTER 5

Conclusions

This chapter summarises the contributions, revisits the research questions, and outlines directions for future work.

5.1. Summary

5.2. Revisiting the Research Questions

5.3. Contributions

5.4. Limitations

5.5. Future Work

References

- Bai, W., Shi, W., de Marvao, A., Dawes, T. J. W., O'Regan, D. P., Cook, S. A., & Rueckert, D. (2015). A bi-ventricular cardiac atlas built from 1000+ high resolution MR images of healthy subjects and an analysis of shape and motion. *Medical Image Analysis*, 26(1), 133–145. <https://doi.org/10.1016/j.media.2015.08.009>
- Bai, W., Sinclair, M., Tarroni, G., Oktay, O., Rajchl, M., Vaillant, G., Lee, A. M., Aung, N., Lubner, E., Nachev, P., et al. (2018). Automated cardiovascular magnetic resonance image analysis with fully convolutional networks. *Journal of Cardiovascular Magnetic Resonance*, 20, 65. <https://doi.org/10.1186/s12968-018-0471-x>
- Beetz, M., Banerjee, A., & Grau, V. (2021). Generating subpopulation-specific biventricular anatomy models using conditional point cloud variational autoencoders. *Statistical Atlases and Computational Models of the Heart (STACOM), MICCAI 2021*. https://doi.org/10.1007/978-3-030-93722-5_4
- Bernard, O., Lalande, A., Zotti, C., Cervenansky, F., Yang, X., Heng, P.-A., Cetin, I., Lekadir, K., Camara, O., et al. (2018). Deep learning techniques for automatic MRI cardiac multi-structures segmentation and diagnosis: Is the problem solved? *IEEE Transactions on Medical Imaging*, 37(11), 2514–2525. <https://doi.org/10.1109/TMI.2018.2837502>
- Campello, V. M., Gkontra, P., Izquierdo, C., Martin-Isla, C., Soez, A., Rodrigo, S. E., et al. (2021). Multi-centre, multi-vendor and multi-disease cardiac segmentation: The M&Ms challenge. *IEEE Transactions on Medical Imaging*, 40(12), 3543–3554. <https://doi.org/10.1109/TMI.2021.3090082>
- Cerqueira, M. D., Weissman, N. J., Dilsizian, V., et al. (2002). Standardized myocardial segmentation and nomenclature for tomographic imaging of the heart. *Circulation*, 105(4), 539–542. <https://doi.org/10.1161/hc0402.102975>
- Chlap, P., Min, H., Vandenberg, N., Dowling, J., Holloway, L., & Haworth, A. (2021). A review of medical image data augmentation techniques for deep learning applications. *Journal of Medical Imaging and Radiation Oncology*, 65(5), 545–563. <https://doi.org/10.1111/1754-9485.13261>
- Choi, H., Moon, G., & Lee, K. M. (2020). Pose2mesh: Graph convolutional network for 3d human pose and mesh recovery from a 2d human pose. *Computer Vision – ECCV 2020*, 769–787. https://doi.org/10.1007/978-3-030-58571-6_45
- Cootes, T. F., Taylor, C. J., Cooper, D. H., & Graham, J. (1995). Active shape models – their training and application. *Computer Vision and Image Understanding*, 61(1), 38–59. <https://doi.org/10.1006/cviu.1995.1004>
- de Marvao, A., Dawes, T. J. W., Shi, W., Minas, C., Rueckert, D., Cook, S. A., & O'Regan, D. P. (2014). Population-based studies of myocardial hypertrophy: High resolution cardiovascular magnetic resonance atlases improve statistical power. *Journal of Cardiovascular Magnetic Resonance*, 16, 16. <https://doi.org/10.1186/s12968-014-0087-0>
- Frangi, A. F., Rueckert, D., Schnabel, J. A., & Niessen, W. J. (2002). Automatic construction of multiple-object three-dimensional statistical shape models: Application to cardiac modeling. *IEEE Transactions on Medical Imaging*, 21(9), 1151–1166. <https://doi.org/10.1109/TMI.2002.804426>

- Garland, M., & Heckbert, P. S. (1997). Surface simplification using quadric error metrics. *Proceedings of the 24th Annual Conference on Computer Graphics and Interactive Techniques*, 209–216. <https://doi.org/10.1145/258734.258849>
- Gropp, A., Yariv, L., Haim, N., Atzmon, M., & Lipman, Y. (2020). Implicit geometric regularization for learning shapes. *Proceedings of the 37th International Conference on Machine Learning*, 3789–3799.
- Grossman, W., et al. (1974). Wall thickness and diastolic properties of the left ventricle. *Circulation*, 49(1), 129–135. <https://doi.org/10.1161/01.CIR.49.1.129>
- Heimann, T., & Meinzer, H.-P. (2009). Statistical shape models for 3d medical image segmentation: A review. *Medical Image Analysis*, 13(4), 543–563. <https://doi.org/10.1016/j.media.2009.05.004>
- Huelnhagen, T., et al. (2017). Myocardial effective transverse relaxation time t_2^* correlates with left ventricular wall thickness: A 7.0 T MRI study. *Magnetic Resonance in Medicine*, 77(6), 2381–2389. <https://doi.org/10.1002/mrm.26312>
- Isensee, F., Jaeger, P. F., Kohl, S. A. A., Petersen, J., & Maier-Hein, K. H. (2021). nnU-Net: A self-configuring method for deep learning-based biomedical image segmentation. *Nature Methods*, 18(2), 203–211. <https://doi.org/10.1038/s41592-020-01008-z>
- Jones, S. E., Buchbinder, B. R., & Aharon, I. (2000). Three-dimensional mapping of cortical thickness using Laplace’s equation. *Human Brain Mapping*, 11(1), 12–32. [https://doi.org/10.1002/1097-0193\(200009\)11:1<12::AID-HBM20>3.0.CO;2-K](https://doi.org/10.1002/1097-0193(200009)11:1<12::AID-HBM20>3.0.CO;2-K)
- Khalafvand, S. S., et al. (2018). Assessment of human left ventricle flow using statistical shape modelling and computational fluid dynamics. *Journal of Biomechanics*, 74, 116–125. <https://doi.org/10.1016/j.jbiomech.2018.04.030>
- Lorensen, W. E., & Cline, H. E. (1987). Marching cubes: A high resolution 3D surface construction algorithm. *Proceedings of the 14th Annual Conference on Computer Graphics and Interactive Techniques (SIGGRAPH)*, 163–169. <https://doi.org/10.1145/37402.37422>
- Martín-Isla, C., et al. (2023). Deep learning segmentation of the right ventricle in cardiac MRI: The M&Ms challenge. *IEEE Journal of Biomedical and Health Informatics*, 27(7), 3302–3313. <https://doi.org/10.1109/JBHI.2023.3267857>
- Medrano-Gracia, P., Cowan, B. R., Ambale-Venkatesh, B., et al. (2014). Left ventricular shape variation in asymptomatic populations: The multi-ethnic study of atherosclerosis. *Journal of Cardiovascular Magnetic Resonance*, 16, 56. <https://doi.org/10.1186/s12968-014-0056-2>
- Mescheder, L., Oechsle, M., Niemeyer, M., Nowozin, S., & Geiger, A. (2019). Occupancy networks: Learning 3d reconstruction in function space. *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 4460–4470. <https://doi.org/10.48550/arXiv.1812.03828>
- Park, J. J., Florence, P., Straub, J., Newcombe, R., & Lovegrove, S. (2019). Deepsdf: Learning continuous signed distance functions for shape representation. *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 165–174. <https://doi.org/10.48550/arXiv.1901.05103>
- Petitjean, C., & Dacher, J.-N. (2011). A review of segmentation methods in short axis cardiac MR images. *Medical Image Analysis*, 15(2), 169–184. <https://doi.org/10.1016/j.media.2010.12.004>
- Qi, C. R., Su, H., Mo, K., & Guibas, L. J. (2017). Pointnet: Deep learning on point sets for 3d classification and segmentation. *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 652–660. <https://doi.org/10.48550/arXiv.1612.00593>
- Ronneberger, O., Fischer, P., & Brox, T. (2015). U-Net: Convolutional networks for biomedical image segmentation. *Medical Image Computing and Computer-Assisted Intervention – MICCAI 2015*, 234–241. https://doi.org/10.1007/978-3-319-24574-4_28

- Shapira, L., Shamir, A., & Cohen-Or, D. (2008). Consistent mesh partitioning and skeletonisation using the shape diameter function. *The Visual Computer*, 24(4), 249–259. <https://doi.org/10.1007/s00371-007-0197-5>
- Shorten, C., & Khoshgoftaar, T. M. (2019). A survey on image data augmentation for deep learning. *Journal of Big Data*, 6, 60. <https://doi.org/10.1186/s40537-019-0197-0>
- Sudlow, C., Gallacher, J., Allen, N., Beral, V., Burton, P., Danesh, J., Downey, P., Elliott, P., Green, J., Landray, M., et al. (2015). UK Biobank: An open access resource for identifying the causes of a wide range of complex diseases of middle and old age. *PLOS Medicine*, 12(3), e1001779. <https://doi.org/10.1371/journal.pmed.1001779>
- Suinesiaputra, A., et al. (2014). A collaborative resource to build consensus for automated left ventricular segmentation of cardiac MR images. *Medical Image Analysis*, 18(1), 50–62. <https://doi.org/10.1016/j.media.2013.09.001>
- Suinesiaputra, A., et al. (2018). Statistical shape modeling of the left ventricle: Myocardial infarct classification challenge. *IEEE Journal of Biomedical and Health Informatics*, 22(2), 503–515. <https://doi.org/10.1109/JBHI.2017.2652449>
- Tancik, M., Srinivasan, P. P., Mildenhall, B., Fridovich-Keil, S., Raghavan, N., Singhal, U., Ramamoorthi, R., Barron, J. T., & Ng, R. (2020). Fourier features let networks learn high frequency functions in low dimensional domains. *Advances in Neural Information Processing Systems*, 33, 7539–7550. <https://doi.org/10.48550/arXiv.2006.10739>
- Taubin, G. (1995). A signal processing approach to fair surface design. *Proceedings of the 22nd Annual Conference on Computer Graphics and Interactive Techniques (SIGGRAPH)*, 351–358. <https://doi.org/10.1145/218380.218473>
- Wang, N., Zhang, Y., Li, Z., Fu, Y., Liu, W., & Jiang, Y.-G. (2018). Pixel2mesh: Generating 3D mesh models from single RGB images. *Computer Vision – ECCV 2018*, 52–67. https://doi.org/10.1007/978-3-030-01252-6_4
- Yezzi, A. J., & Prince, J. L. (2003). An Eulerian PDE approach for computing tissue thickness. *IEEE Transactions on Medical Imaging*, 22(10), 1332–1339. <https://doi.org/10.1109/TMI.2003.817775>

APPENDIX A

Appendix Title

A.1. Section in the Appendix